


Reinforcement Learning (RL)

Symbols

a action, arm (multi-armed bandit problem)	m momentum
A advantage	Pr transition distribution
α learning rate, step size	π policy (<i>state</i> \rightarrow <i>action</i>)
b bandit, baseline estimate, bias	q quality
c cost, context (= state)	r reward
d difference error	R return
Δ difference	ρ regret
D replay buffer	s state
E eligibility trace	t time
ϵ exploration rate (0 1)	τ trajectory
ϕ state transition function	θ parameter weight (learning target)
g gain (reward over time)	v value
γ discount factor	vm variance momentum
∇ gradient (spacial derivative)	w weight, winning probability
H horizon	y target / prediction
J expected return (see loss)	$*$ optimal
L expected (squared) loss (see return)	$'$ next / derivative
λ trace decay (0 1)	$\hat{\cdot}$ estimation
	$ $ vector norm

Agent-Environment Interface



The Agent at each step t receives a representation of the environment's *state*, $S_t \in S$ and it selects an action $A_t \in A(s)$. From its action the agent receives a *reward*, $R_{t+1} \in R \in \mathbb{R}$.

Bandits

- solve multi-armed bandit problem
- no state
- reduces RL to exploration vs exploitation

They estimate $\hat{r}_a \approx \mathbb{E}[r|a]$

Greedy

$$\begin{aligned} n_a &= \sum_t 1 \\ \hat{r}_a &= \sum_t \frac{r_t}{n_a} \\ a_t &= \operatorname{argmax}_{a \in A} \hat{r}_a \end{aligned} \quad (1)$$

Optimistic-Greedy

large initial \hat{r}_a

ϵ -Greedy

for probability ϵ , pick randomly

Upper Confidence Bound (UCB)

try all for k rounds, then

$$a_t = \operatorname{argmax}_{a \in A} \hat{r}_a + \sqrt{\frac{2 \log t}{n_a}} \quad (2)$$

Contextual Bandit

Linear UCB (LinUCB)

$$\begin{aligned} x_{t,a} &= \Phi(s_{t,a}) \\ \mathbb{E}[r_{t,a}|x_{t,a}] &= \theta_a \cdot x_{t,a} \\ \hat{\theta}_a &= A^{-1}b \end{aligned} \quad (3)$$

Posterior / Thompson sampling

Pick actions by probability they maximize expected reward

Greedy in the Limit with Infinite Exploration (GLIE)

- infinitely explores
- converges on greedy

Concepts

Reward

$$G_t = \sum_{k=0}^H \gamma^k r_{t+k+1} \quad (4)$$

for *discount factor* γ and (infinite?) *horizon* H .

Policy

A *policy* maps a state to an action

$$\pi_t(a|s) \quad (5)$$

probability to pick an action $A_t = a$ if $S_t = s$.

Markov Decision Process (MDP)

a 5-tuple (S, A, P, R, γ)

$$\begin{aligned} \text{state transition probabilities:} \\ p(s'|s, a) &= Pr\{S_{t+1} = s' | S_t = s, A_t = a\} \\ \text{expected reward for state-action-nextstate:} \\ r(s', s, a) &= \mathbb{E}[R_{t+1} | S_{t+1} = s', S_t = s, A_t = a] \end{aligned} \quad (6)$$

Value Function

Value (expected return, total discounted reward) of a specific state s under policy π for a MDP:

$$\begin{aligned} v_\pi(s) &= \mathbb{E}_\pi[G_t | S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a)[r + \gamma v_\pi(s')] \\ v_*(s) &= \max_\pi v_\pi(s) \end{aligned} \quad (7)$$

Action-Value (Q) Function

expected reward for state-action pairs:

$$\begin{aligned} q_\pi(s, a) &= \mathbb{E}_\pi[G_t | S_t = s, A_t = a] \\ &= \sum_{s', r} p(s', r|s, a)[r + \gamma V_\pi(s')] \\ q_*(s, a) &= \max_\pi q_\pi(s, a) \end{aligned} \quad (8)$$

rewriting v_* for $q_*(s, a)$:

$$v_*(s) = \max_{a \in A(s)} q_{\pi_*}(s, a) \quad (9)$$

i.e. state value under optimal policy = expected return from its best action.

Bellman Equation

Recursive property *Value* 7 / *Q* 8 functions

Contraction Mapping

For metric space (X, d) and $f : X \rightarrow X$, f is a *contraction* given a *Lipschitz coefficient* $k \in [0, 1)$ where for all x / y in X :

$$d(f(x), f(y)) \leq kd(x, y) \quad (10)$$

Contraction Mapping theorem

For complete metric space (X, d) and contraction $f : X \rightarrow X$, there is only 1 fixed point x^* where $f(x^*) = x^*$.

For point x in X , and $f^n(x)$ inductively defined by $f^n(x) = f(f^{n-1}(x))$, $f^n(x) \rightarrow x^*$ as $n \rightarrow \infty$, yielding a unique optimal solution for DP.

Model-based Methods (known MDP)

Exhaustive Search

brute force, usually computationally unviable.

Dynamic Programming (DP)

+ bootstrap (learn mid-episode)

Find π_* for V / Q :

Policy Iteration

Initialize $V(s) \in \mathbb{R}$.e.g. 0, $\Delta \leftarrow 0$, $\pi(s) \in A$ for all $s \in S$

1. Policy Evaluation

while $\Delta < \theta$ (e.g. 0.001) **do**

foreach $s \in S$ **do**

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a)[r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

end

end

2. Policy Improvement

policy-stable \leftarrow *true*

while not *policy-stable* **do**

foreach $s \in S$ **do**

$old_action \leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s', r} p(s', r|s, a)[r + \gamma V(s')]$

policy-stable \leftarrow *old-action* \neq $\pi(s)$

end

end

Policy iteration methods:

- gradient-based (policy gradient methods): gradient ascent
- gradient-free: simulated annealing, cross-entropy search or methods of evolutionary computation
- value search/iteration: stop after 1 state sweep
- async DP: update iteratively, no full sweeps
- generated policy iteration (GPI)

Value Iteration

ditch $V(s)$ convergence for policy improvement and truncated policy eval step in 1 operation:

```
Initialize  $V(s) \in \mathbb{R}$ , e.g. 0,  $\Delta \leftarrow 0$ 
while  $\Delta < \theta$  (e.g. 0.001) do
  foreach  $s \in S$  do
     $v \leftarrow V(s)$ 
     $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$ 
     $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
  end
end
output: deterministic policy  $\pi \approx \pi_*$  where
 $\pi(s) = \arg\max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$ 
```

Model-free methods

Monte Carlo (MC) Methods

- uses **averaging sample returns** per state-action pair
- episodic
- + sampling

```
Initialize for all  $s \in S, a \in A(s)$  :
 $Q(s,a) \leftarrow$  arbitrary
 $\pi(s) \leftarrow$  arbitrary
 $Returns(s,a) \leftarrow$  empty list
while forever do
  Pick  $S_0 \in S$  and  $A_0 \in A(S_0)$ , all  $p(s,a) > 0$ 
  Generate an episode starting at  $S_0, A_0$  following  $\pi$ 
  foreach pair  $s,a$  in the episode do
     $G \leftarrow$  return for first occurrence of  $s,a$ 
    Append  $G$  to  $Returns(s,a)$ 
     $Q(s,a) \leftarrow average(Returns(s,a))$ 
  end
  foreach  $s$  in the episode do
     $\pi(s) \leftarrow \arg\max_a Q(s,a)$ 
  end
end
```

estimate for non-stationary problems:

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)] \quad (11)$$

for learning rate α , how much we want to forget about past experiences.

Temporal Difference (TD)

- + DP's bootstrap
- + MC's sampling
- substitutes expected discounted reward G_t from the episode with an estimation:

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] \quad (12)$$

State-action-reward-state-action (SARSA)

- on-policy TD control
- can use priors

```
Initialize  $Q(s,a)$  arbitrarily and
 $Q(\text{terminal} - \text{state},) = 0$ 
foreach episode  $\in episodes$  do
  Pick  $a$  from  $s$  by policy from  $Q$  (e.g.  $\epsilon$ -greedy)
  while  $s$  is not terminal do
    Take action  $a$ , observe  $r, s'$ 
    Pick  $a'$  from  $s'$  by policy from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $y \leftarrow r + \gamma Q(s', a')$ 
     $Q(s,a) \leftarrow Q(s,a) + \alpha[y - Q(s,a)]$ 
     $s \leftarrow s'$ 
     $a \leftarrow a'$ 
  end
end
```

n -step Sarsa (n -step TD) for n -step Q-Return:

$$\begin{aligned} q_t^{(n)} &= \gamma^n Q(S_{t+n}) + \sum_{i=1}^n \gamma^{i-1} R_{t+i} \\ Q(s_t, a_t) &\leftarrow Q(s_t, a_t) + \alpha [q_t^{(n)} - Q(s_t, a_t)] \end{aligned} \quad (13)$$

Forward View Sarsa(λ) (/ TD(λ))

$$\begin{aligned} q_t^\lambda &= (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)} \\ Q(s_t, a_t) &\leftarrow Q(s_t, a_t) + \alpha [q_t^\lambda - Q(s_t, a_t)] \end{aligned} \quad (14)$$

Backward View Sarsa(λ) (/ TD(λ))

- + more efficient
- + can update at every time-step
- eligibility traces : find cause in frequency vs recency

$$\begin{aligned} E_0(s,a) &= 0 \\ E_t(s,a) &= \gamma \lambda E_{t-1}(s,a) + \mathbf{1}(S_t = t, A_t = a) \\ \delta_t &= R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \\ Q(s,a) &\leftarrow Q(s,a) + \alpha \delta_t E_t(s,a) \end{aligned} \quad (15)$$

Linear Function Approximation

- + efficient
- + generalize

update temporal difference error, minimize squared loss:

$$\begin{aligned} \delta &\leftarrow r_t + \gamma \theta^T \phi(s_{t+1}) - \theta^T \phi(s_t) \\ \theta &\leftarrow \theta + \alpha \delta \phi(s_t) \\ J(\theta) &= ||\delta||^2 \end{aligned} \quad (16)$$

Q Learning

$$\begin{aligned} \delta &\leftarrow r_t + \gamma \arg\max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-) - Q(\phi(s_t, a); \theta_i) \\ Q(s,a) &\leftarrow Q(s,a) + \alpha [r + \gamma \arg\max_{a'} Q(s', a') - Q(s,a)] \end{aligned} \quad (17)$$

Replay Memory

update θ by SGD

$$\nabla_{\theta_i} J(\theta_i) = \mathbb{E}_{(s_t, a_t, r_t, s_{t+1}) \sim D} \delta \nabla_{\theta_i} Q(\phi(s_t, a_t); \theta) \quad (18)$$

Deep Q Learning (DQL)

Made by *DeepMind*, uses a deep neural net (Q -network) for the Q function. Keeps N observations in a *memory* to train on.

$$\begin{aligned} y &= r_t + \gamma \arg\max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-) \\ \nabla_i J(\theta_i) &= \mathbb{E}_{(s,a,r,s') \sim U(D)} [(y - \underbrace{Q(s,a;\theta_i)}_{\text{prediction}})^2] \end{aligned} \quad (19)$$

for network weights θ and experience replay history $U(D)$.

Initialize replay memory D with capacity N

```
Initialize  $Q(s,a)$  arbitrarily
foreach episode  $\in episodes$  do
  Pick  $a$  from  $s$  by policy from  $Q$  (e.g.  $\epsilon$ -greedy)
  while  $s$  is not terminal do
    Take action  $a$ , observer  $r, s'$ 
    Store transition  $(s,a,r,s')$  in  $D$ 
    Sample random transitions from  $D$ 
     $y_i \leftarrow \begin{cases} r_j & \text{for terminal } s'_j \\ r_j + \gamma \max_a Q(s', a'; \theta) & \text{otherwise} \end{cases}$ 
    Perform gradient descent step on
     $(y_j - Q(s_j, a_j; \Theta))^2$ 
     $s \leftarrow s'$ 
  end
end
```

Prioritized Replay Memory learn esp. from high loss (traumas)

$$p(s_t, a_t, r_t, s_{t+1}) \propto r_t + \gamma \max_{a \in A} Q(\Phi(s_{t+1}, a); \theta_i^-) \quad (20)$$

Double DQN

solve bias from reused θ , faster

$$y = r_t + \gamma Q(\Phi(s_{t+1}, \arg\max_{a \in A} Q(\Phi(s_{t+1}, a); \theta_i^-)); \theta_i^-) \quad (21)$$

Direct Policy Search

learn π .

Policy Gradient

- + simpler than Q or V
- + allows stochastic policy (rock-paper-scissors)
- local optimum
- less sample efficient

$$L = \mu(\log p(a|s) \cdot R(s)) \quad (22)$$

Likelihood Ratio

return state-action trajectory:

$$R(\tau) = \sum_{t=0}^T R(s_t, a_t) \quad (23)$$

expected return:

$$\begin{aligned} J(\theta) &= \mathbb{E}[\sum_{t=0}^T R(s_t, a_t); \pi_\theta] \\ &= \sum_{\tau} R(s_t, a_t); \pi_\theta \end{aligned} \quad (24)$$

find θ to max:

$$J(\theta) = \sum_{\tau} P(\tau; \theta) R(\tau) \quad (25)$$

yielding:

$$\max_{\theta} J(\theta) = \max_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) \quad (26)$$

$$\nabla_{\theta} J(\theta) = \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau) \quad (27)$$

sampled over m trajectories:

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m P(\tau; \theta) \nabla_{\theta} \log P(\tau^i; \theta) R(\tau^i) \quad (28)$$

gradient chases reward.

REINFORCE

- Get info on states/actions/rewards for policy during episode
- Calc episode return for collected rewards
- Update model params toward policy gradient

$$\nabla_{\theta} J(\theta) = \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R \quad (29)$$

desired loss function:

$$\frac{1}{m} \sum_{t=1}^m \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R \quad (30)$$

Baselined REINFORCE

use baselined rewards

$$\nabla_{\theta} J(\theta) = \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (R - V_{\phi}(s_t)) \quad (31)$$

Actor-critic

+ less variance than Baselined REINFORCE

- actor (makes policy): policy gradient
- critic: policy iteration

advantage:

$$A_{\pi}(s, a) = Q(s, a) - V(s) \quad (32)$$

Asynchronous Advantage Actor Critic (A3C)

- 5-step Q-Value estimation
- shares params between actor/critic

+ run parallel, one policy

+ no more need for DQN's replay policy

github.com/tycho01/Reinforcement-Learning-Cheat-Sheet